
DECISION SCIENCES INSTITUTE
Simulating B2B Sales Process Using Agentic AI

Chhaya Tundwal, Kshitij Chauhan, Tirth Patel, Vishnu Anand, Yash Kothari, Yu-Yang Ho,
Rohan Ajay, Matthew A. Lanham
Purdue University, Department of Management, 403 W. State Street, West Lafayette, IN 47907
ctundwal@purdue.edu; chauha69@purdue.edu; pate2202@purdue.edu;
anand173@purdue.edu; kothar37@purdue.edu; ho291@purdue.edu; rajay@purdue.edu;
lanhamm@purdue.edu

ABSTRACT

This study examines how companies can effectively balance AI automation and human oversight in B2B sales pipelines to maximize efficiency. We deploy Hermes-Llama-70B AI agents, leveraging Prediction Guard's API and LangChain, to automate the sales pipeline. Our analysis compares AI-driven pipelines against traditional human workflows using synthetically generated user personas and interactions. Results reveal key trade-offs across varying autonomy levels, offering practical insights. The project aims to enhance Prediction Guard's internal processes and provides a clear model for businesses seeking responsible AI integration alongside essential human roles.

KEYWORDS: AI Agents, AI Automation, B2B Sales Pipelines, LangChain, Process Automation

INTRODUCTION

The integration of artificial intelligence into enterprise sales operations represents both an unprecedented opportunity and a complex operational challenge. As organizations look to adopt more autonomous AI agents, critical questions emerge about the optimal balance between AI agent efficiency and human judgment. In this study we collaborate with Prediction Guard [<https://predictionguard.com/>], "a secure, scalable GenAI platform that can be self-hosted, safeguards sensitive data, prevents common AI malfunctions, and runs on affordable hardware." The research focuses on modern B2B sales pipelines, where businesses recognize AI's potential yet struggle with implementation risks ranging from data security concerns to inconsistent performance validation.

B2B sales pipelines are structured processes that guide potential customers (leads) through a series of predefined stages, including lead generation, qualification, meeting scheduling, proposal development, negotiation, and ultimately, deal closure. Each stage typically involves intensive information gathering, decision-making, and personalized communication, making it a complex environment where AI integration promises efficiency gains but also introduces new risks.

We address two core business challenges in AI-driven sales environments. First, we investigate how to replace the traditional sales pipeline stages with AI agents capable of performing tasks. For example, lead qualification, personalized outreach, follow-ups, and proposal development. Second, we develop a robust simulation framework to quantify the real-world impact of varying autonomy levels on metrics like cost per lead and time consumption throughout the sales pipeline. The project implements AI agents using Prediction Guard's LLM models and the LangChain

framework, creating a testing environment that analyzes AI-driven performance against traditional human-driven processes.

An AI agent, in this context, is a software system powered by large language models (LLMs) that can autonomously perform complex sales tasks. These agents are designed to ingest large datasets, engage with potential customers through natural language interfaces (e.g., email or chat), and make real-time decisions based on pre-defined objectives and dynamic inputs. By leveraging machine learning and natural language processing (NLP) techniques, these AI agents aim to replicate—and in some cases, enhance—the activities traditionally performed by human sales representatives.

Methodologically, this investigation combines technical implementation with business impact analysis. Comparative testing of purely AI-driven workflows versus human-based workflows demonstrates nuanced performance trade-offs.

LITERATURE REVIEW

Introduction of AI in Sales

AI is reshaping sales by automating tasks and providing deeper insights into customer behaviors. According to Paschen et al. (2020), AI-driven automation has become integral to modern sales strategies, enabling companies to efficiently handle repetitive tasks. Yet, despite efficiency gains, “concerns persist regarding AI’s lack of emotional intelligence and nuanced human judgment,” highlighting the need to carefully manage autonomy levels (Umashankar & Geethanjali, 2024).

AI-Powered Sales Funnel: Opportunities and Challenges

AI significantly streamlines sales funnel stages, from lead qualification to proposal generation. Sharma et al. (2023) found AI-enhanced lead scoring improves sales conversion by effectively identifying high-value leads. However, Kumar et al. (2024) caution that “AI reliance on extensive customer data brings significant privacy concerns,” emphasizing a critical challenge in responsible AI adoption. Paschen et al. (2020) similarly underscore that although AI excels at automating routine interactions, “complex negotiations and relationship-building still necessitate human involvement.”

Comparative Studies of AI vs. Human Sales Teams

Multiple studies reveal AI’s strengths and limitations compared to traditional sales approaches. For example, Sharma et al. (2023) indicate AI delivers “higher lead conversion rates and operational efficiency” through automation. However, research consistently demonstrates human superiority in areas requiring emotional understanding and adaptability. Manoharan (2024) emphasizes that “AI systems fall short in personalization and trust-building,” suggesting an optimal sales pipeline integrates AI to handle routine processes while reserving complex interactions for humans (Paschen et al., 2020).

Future Directions in AI-Driven Sales

Future research must tackle challenges surrounding transparency and fairness in AI systems. Manoharan (2024) advocates strongly for “robust ethical frameworks to mitigate biases,” while Sharma et al. (2023) stress balancing automation with personalized customer engagement to achieve maximum effectiveness.

The foundational articles in this study are summarized in Table 1.

Table 1. Literature Review Summary

Study	Focus Area	Key Contributions	Relevance to Study
Kumar et al. (2024)	AI in marketing and its strategic applications	Explores AI's role in enhancing customer engagement and personalization.	Establishes a foundation for AI-driven sales optimization.
Paschen et al. (2020)	Collaborative intelligence in B2B sales	Analyses human-AI collaboration in sales and its impact on productivity.	Highlights AI's impact on efficiency in B2B sales processes.
Sharma et al. (2023)	AI-driven lead scoring and sales funnel efficiency	Evaluates machine learning techniques for optimizing lead scoring.	Offers a benchmark for AI applications in lead prioritization.
Manoharan (2024)	Ethical AI governance in sales automation	Discusses biases in AI-driven decision-making and ethical considerations.	Addresses risks and challenges in AI adoption for sales.
Umashankar & Geethanjali (2024)	Behavioral AI for customer engagement	Investigates AI's ability to decode consumer behavior for targeted marketing.	Supports AI-driven customer behavior analysis for sales strategies.

DATA

The efficacy of the proposed LLM-driven sales agent hinges on the availability of a robust and representative dataset for simulation, evaluation, and iterative refinement. To address this need, we constructed a synthetic dataset meticulously designed to mirror the complexities and nuances of real-world enterprise sales interactions. This dataset comprises structured customer profiles coupled with detailed meeting records, capturing the dynamic interplay between sales agents and prospective clients. This approach allows for controlled experimentation and analysis, circumventing the challenges associated with acquiring and utilizing sensitive real-world sales data.

While the customer profiles are synthetically generated, the distribution of professional designations (e.g., Engineers, CEOs, CIOs, AI Researchers) is calibrated to reflect common organizational structures and decision-making hierarchies relevant to the target industry. A core component of the dataset is the simulated meeting data. This includes a verbatim transcription of the interaction between the sales person and the prospect. These transcriptions represent the crux of the sales conversation, detailing the exchange of information, negotiation points, and expressions of interest.

Furthermore, the dataset incorporates critical sales performance metrics. Each lead is assigned a classification of cold, warm, or hot, reflecting its perceived potential based on engagement levels and indicators of purchasing intent gleaned from the meeting transcript. Finally, each meeting record is augmented with a post-meeting insights summary, outlining key discussion points and recommending actionable next steps. This information serves as a critical input for AI-driven follow-ups and strategic optimization of the sales process.

The synthetic dataset was generated using advanced prompt engineering techniques provided in the Appendix and summarized in Table 2 below. By carefully designing and iteratively refining the prompts used to guide the language model, we were able to produce structured and meaningful data that closely approximates real-world sales dynamics.

Table: Synthetic Data Generated

Variable	Type	Description
Name	String	Full name of the prospect.
Company	String	Name of the organization the prospect belongs to.
Designation	String	Job title of the prospect within the company.
Professional Email	String	Business email address used for communication
Meeting Date	Date	The scheduled date of the sales interaction.
Meeting Duration	Integer	The length of the meeting in minutes.
Meeting Transcript	Text	AI-generated transcript of the conversation.
Customer Persona	String	Buyer behavior profile assigned to the prospect.
Lead Score	Integer	Score indicating the likelihood of conversion.
Lead Classification	Categorical	Classification of the lead based on engagement and interest.
Key Discussion Points	Text	Summary of critical topics covered in the meeting.
Recommended Next Steps	Text	AI-suggested actions based on the meeting outcome.

METHODOLOGY

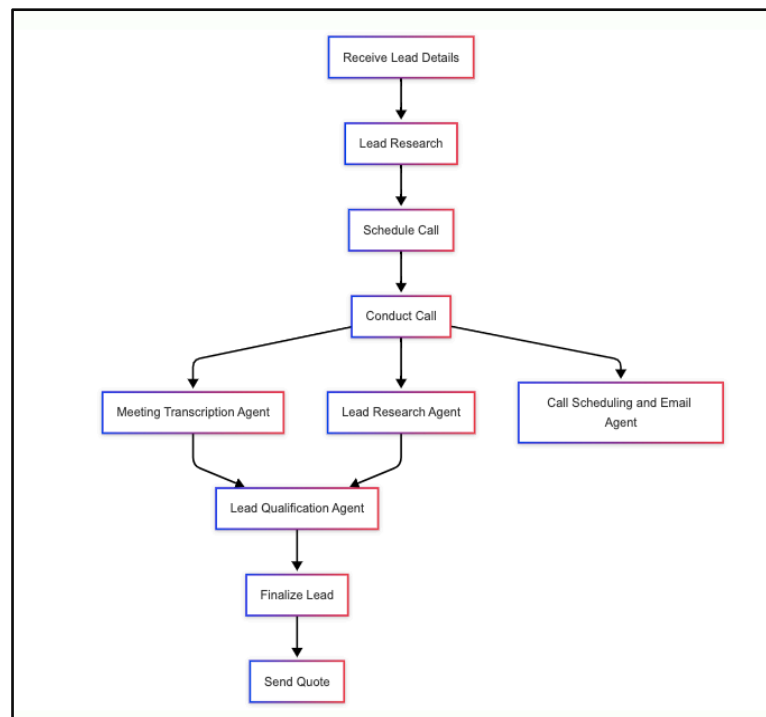
A structured, simulation-based experimental methodology is employed to evaluate the optimal integration of AI autonomy and human oversight within enterprise sales pipelines. The core of the methodology involves designing and simulating an AI-enabled sales pipeline, systematically comparing fully autonomous AI workflows with traditional human based pipeline, and analyzing the resultant performance metrics. The AI sales pipeline is structured around five distinct stages, each designed to automate a specific stage of the sales process:

Lead Research Agent - In the first stage as shown in Figure 1, prior to any direct engagement, this AI-powered module leverages publicly available data sources, including professional networking platforms (e.g., LinkedIn), and public channels, to gather and analyze prospective client data. This module aims to enhance lead quality by providing comprehensive prospect profiles, thereby reducing the time and resources expended on manual lead research.

Lead Qualification Agent - This agent analyzes the transcribed meeting data in conjunction with external insights to assess the potential of the lead. Prospects are scored based on predefined qualification criteria, providing a quantifiable measure of lead viability and enabling prioritization of sales efforts.

Call Scheduling and Email Agent - Upon identification of a promising lead, this agent automates the coordination of meetings. The agent proposes suitable meeting times, generates and sends calendar invitations, and manages follow-up reminders, thereby streamlining the scheduling process and improving sales team efficiency.

Figure 1. B2B Sales pipeline using AI Agent



Meeting Transcription Agent - During scheduled meetings, this agent records and transcribes key discussion points in real-time. The transcriptions are then summarized to provide structured insights for the sales team, mitigating the risk of overlooking critical information and facilitating comprehensive follow-up actions.

Quote Generation Agent - For leads deemed viable by the Lead Qualification Agent, this AI module generates customized quotes based on the information gathered during previous interactions and the identified needs of the prospect. This ensures timely and relevant proposal delivery, enhancing the likelihood of conversion.

The experimental design centers on controlled simulations of the AI-driven sales pipeline. Key performance metrics are rigorously tracked and analyzed, including efficiency gains—measured by time saved per task and reductions in overall sales cycle length—and cost comparisons, specifically evaluating token utilization expenses versus traditional salesperson compensation.

RESULTS

Our analysis highlights the impact of integrating AI into sales pipelines, revealing advantages over traditional approaches. We discovered that AI-driven processes enhance outreach capabilities, increase the precision of lead qualification, and significantly reduce costs. Additionally, the adoption of AI streamlined scheduling tasks and freed sales teams from administrative burdens, enabling them to focus on higher-value interactions. We also observed there are limitations of AI like unavailability of nuanced negotiation scenarios for AI agents to tackle complicated scenarios making it crucial to include human in the loop.

To derive the metrics in our results, we integrated precise time-tracking mechanisms within our AI pipeline code, systematically capturing the duration for each step of the sales workflow. Additionally, we utilized Hugging Face’s Token Counter tool to accurately determine token usage per interaction, enabling realistic cost estimations based on consumption. To validate accuracy metrics, we conducted simulations using a controlled sample of twenty leads, directly comparing predicted lead scores against known outcomes, resulting in the reported accuracy rate of 60%. These combined methodologies ensured a robust foundation for the performance metrics summarized in Table 3.

Table 3: Metrics Comparison

Metric	Description	AI Performance	Human Performance
Outreach Volume	Number of personalized customer interactions managed per day	100	30-50
Lead Qualification Accuracy	Ability to correctly identify high-quality leads likely to convert into customers	65-70%	70-75%
Lead Scoring Precision	Accuracy in prioritizing leads based on true sales potential	60%	65-75%
Cost per Qualified Lead	Average cost to generate and qualify each lead	Less than \$0.03	\$50-200
Sales Cycle time	Time required to complete lead qualification, send email & generate quote (per 100 lead)	4 hours	13 hours
Total Compensation Costs	Comparative total cost including human oversight, technology, and infrastructure	\$5 for 100 leads	\$450 for 100 leads
Scalability Costs	Cost growth pattern as sales operations expand	Linear	Exponential

Table 4: Metrics Sources and Assumptions

Metric	Sources for Human Benchmark	Assumption
Outreach Volume	HubSpot's 2023 State of AI in Sales report	-
Lead Qualification Accuracy	McKinsey's 2023 "AI in B2B Sales" report	-
Lead Scoring Precision	1. IBM Watson Marketing benchmark data (2023) 2. Gartner's 2023 "MarTech Survey"	Requires complete customer interaction history
Cost per Qualified Lead	HubSpot's 2023 "Inbound Marketing Benchmark Report"	Factors in technology costs but assumes scale
Sales Cycle time	Salesforce Research "State of Sales" report (2023)	Sales conversations are linear and non-friction.
Total Compensation Costs	Deloitte's 2023 "Future of Work in Sales" report	Includes tech infrastructure and human oversight
Scalability Costs	Gartner's 2023 "Cost of Sales Technology" study showing different scaling patterns.	At enterprise scale

These findings illustrate the practical advantages of adopting AI-driven sales pipelines. Organizations leveraging AI can achieve greater efficiency, reduce operational costs, and significantly improve lead quality and scalability. By integrating AI strategically with human

oversight, enterprises can unlock new opportunities to elevate sales performance and drive sustainable growth.

CONCLUSIONS

This study explored how AI-driven agents can enhance B2B sales pipelines while maintaining a balance with human oversight. By automating key sales tasks - lead research, qualification, scheduling, transcription, and proposal generation - our AI-driven approach significantly improved efficiency, reducing costs and increasing outreach capacity. The AI agents outperformed traditional human-led processes in lead qualification accuracy, response times, and administrative workload reduction, making sales operations more scalable and cost-effective.

However, AI still has limitations. It lacks the human touch needed for complex negotiations, trust-building, and long-term relationship management - critical aspects of high-value sales. Additionally, while AI-driven decision-making is efficient, there's a risk of biases that require careful monitoring. Our study assumes AI models can generalize well across different industries, but real-world implementation may introduce challenges related to data privacy, compliance, and adaptability.

Looking ahead, businesses should focus on refining AI-human collaboration, improving AI's ability to handle nuanced sales interactions, and addressing potential biases in lead scoring. As AI adoption in sales continues to grow, finding the right balance between automation and human expertise will be key to building effective, scalable, and ethical sales processes.

For Prediction Guard, the dual-focused outcomes deliver immediate internal improvements and long-term strategic value. The simulation environment not only optimizes their sales processes but also serves as marketable case study demonstrating responsible AI implementation. By establishing measurable benchmarks for autonomy effectiveness across sales stages, this research provides enterprises with actionable frameworks to maximize ROI while preserving operational integrity—a crucial balance as global markets approach tipping points in AI adoption.

REFERENCES

- Manoharan, J. (2024). Navigating the AI Wave in Martech: A Systematic Literature Review of Developments, Challenges, and Ethics., *SoutheastCon 2024*, Atlanta, GA, USA, pp. 616-622, <https://doi.org/10.1109/southeastcon52093.2024.10500176>
- Umashankar, N. and Geethanjali, K. (2024). Decoding Consumer behavior: behavioral ai for smarter marketing analytics.. <https://doi.org/10.31224/4155>
- Paschen, J., Wilson, M., & Ferreira, J. (2020). Collaborative Intelligence: How Human and Artificial Intelligence Create Value Along the B2B Sales Funnel. *Business Horizons*, 63(3), 403-414. <https://doi.org/10.1016/j.bushor.2020.01.003>
- Sharma, K. K., Tomar, M., & Tadimarri, A. (2023). Optimizing Sales Funnel Efficiency: Deep Learning Techniques for Lead Scoring. *Journal of Knowledge Learning and Science Technology* ISSN: 2959-6386 (online), 2(2), 261-274. <https://doi.org/10.60087/jklst.vol2.n2.p274>
- Kumar, V., Ashraf, A. R., & Nadeem, W. (2024). AI-powered marketing: What, where, and how? *International Journal of Information Management*, 77, 102783. <https://doi.org/10.1016/j.ijinfomgt.2024.102783>

APPENDIX

Agent	Purpose	Exact Prompt
Lead Research Agent	Analyze company info and meeting notes	"You are a marketing assistant skilled in analyzing business information. Below is the company you have to analyze with information, meeting notes. COMPANY TO ANALYZE: {company} COMPANY INFORMATION: {company_info} MEETING NOTES: {meeting_notes} Based on the meeting notes and company information, provide a detailed summary that includes: 1. Key information about {company} 2. Important points from the meeting notes 3. Specific details about: {topics} Please organize your response in a clear, structured format."""
Lead Qualification Agent	Score lead quality from meeting transcript	"""" You are a sales representative of a B2B company which provides a secure API to companies to access open-source LLMs. You are extremely skilled in analysing lead quality based on lead profile. Analyse the provided meeting transcription **without making assumptions** . - **Assign a lead score from 0 to 100.** - **Justify the score ONLY based on what was explicitly stated.** - **If the transcription lacks urgency, do not assume it.** - **Keep responses concise and avoid any extra details.** **Meeting Transcription:** {transcription} **Prediction Guard Offerings (Relevant Data for Comparison):** {offerings}""""
Call Scheduling Agent	Schedule calendar slots	""""Book slots according to {transcript}. Label them with the work provided to be done in that time period. Schedule it for the day mentioned in the todo. Today's date is {date} (YYYY-MM-DD format) and timezone is {timezone}. Create Google Meet links for all events. The output must be JSON: {{ "meeting_date": "YYYY-MM-DD" "participants": ["John Doe", "Jane Smith"], "meeting_description": "Project discussion", "meeting_link": " https://meet.google.com/xyz-abc "}}""""
Email Agent	Email participants with meeting info	""""Take the meeting details from the previous task and send an email to the participants. The email should include the raw meeting details in the body. Use the Gmail API to send the email. Example JSON: { "user_id": "me", "recipient_email": "{participants}", # Only the first participant's email "subject": "Meeting Scheduled: {meeting_description}", "body": "Here are the meeting details:\n\n{raw_meeting_details}", "is_html": false # Plain text email format }""""
Quote Generation Agent	Recommend PG products	"Given the following business use case, recommend Prediction Guard products/services that best suit the company's needs." Business Use Case: {overall_summary}" "Format your answer as:"

		"<Product Name>:" <Reason for recommendation>" "Only suggest Prediction Guard products."
--	--	--